

1. Air Quality Analysis: Inbuilt dataset: seaborn.load_dataset('mpg') in Python

A. Analyze missing values in the dataset and impute them appropriately.

B. Find the average ozone levels per month

```
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

df = sns.load_dataset('mpg')

print(df.isnull().sum())
df['horsepower'] = df['horsepower'].fillna(df['horsepower'].mean())

avg_mpg_by_year = df.groupby('model_year')['mpg'].mean().reset_index()
print(avg_mpg_by_year)
```

mpg	0
cylinders	0
displacement	0
horsepower	6
weight	0
acceleration	0
model_year	0
origin	0
name	0

dtype: int64

	model_year	mpg
0	70	17.689655
1	71	21.250000
2	72	18.714286
3	73	17.100000
4	74	22.703704
5	75	20.266667
6	76	21.573529
7	77	23.375000
8	78	24.061111
9	79	25.093103
10	80	33.696552
11	81	30.334483
12	82	31.709677

1. Car Performance Analysis: Inbuilt dataset: seaborn.load_dataset('mpg')

- Display the first 5 rows of the dataset.
- How many rows and columns does the dataset have?
- What are the names of all the columns in the dataset?
- Find the average miles per gallon (mpg) for each number of cylinders.

- Create a scatter plot to show the relationship between horsepower and mpg.

```
# dataset
df = sns.load_dataset('mpg')

print(df.head())

print(df.shape) # (rows, columns)

print(df.columns.tolist())


avg_mpg_by_cyl = df.groupby('cylinders')['mpg'].mean()
print(avg_mpg_by_cyl)

sns.scatterplot(data=df, x='horsepower', y='mpg')
plt.title('Horsepower vs MPG')
plt.xlabel('Horsepower')
plt.ylabel('Miles Per Gallon (MPG)')
plt.show()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	\
0	18.0	8	307.0	130.0	3504	12.0	
1	15.0	8	350.0	165.0	3693	11.5	
2	18.0	8	318.0	150.0	3436	11.0	
3	16.0	8	304.0	150.0	3433	12.0	
4	17.0	8	302.0	140.0	3449	10.5	

	model_year	origin	name
0	70	usa	chevrolet chevelle malibu
1	70	usa	buick skylark 320
2	70	usa	plymouth satellite
3	70	usa	amc rebel sst
4	70	usa	ford torino

```
(398, 9)
['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
'acceleration', 'model_year', 'origin', 'name']
cylinders
3    20.550000
4    29.286765
5    27.366667
6    19.985714
8    14.963107
Name: mpg, dtype: float64
```



1. Titanic Survival Analysis: Inbuilt Dataset: `seaborn.load_dataset('titanic')` in Python

A. Compute the survival rate grouped by gender (sex) and passenger class (class).

B. Filter and display records of passengers who:

- Were in 1st class,
- Are female, and
- Had a fare greater than 50.

```
df = sns.load_dataset('titanic')

survival_rate = df.groupby(['sex', 'class'])
['survived'].mean().reset_index();
print(survival_rate)

filtered_passengers = df[(df['sex'] == 'female') &
                          (df['class'] == 'First') &
                          (df['fare'] > 50)]

print(filtered_passengers[['sex', 'class', 'fare']]);
```

	sex	class	survived
0	female	First	0.968085

```

1  female  Second  0.921053
2  female  Third  0.500000
3   male   First  0.368852
4   male   Second 0.157407
5   male   Third  0.135447
   sex    class    fare
1  female  First    71.2833
3  female  First    53.1000
31 female  First   146.5208
52 female  First    76.7292
61 female  First    80.0000
..      ...      ...
835 female  First    83.1583
849 female  First    89.1042
856 female  First   164.8667
871 female  First    52.5542
879 female  First    83.1583

```

```
[82 rows x 3 columns]
```

```

/var/folders/kx/41v1tt6jlyx79h_8wk_hkl8w0000gn/T/
ipykernel_4467/1975499764.py:4: FutureWarning: The default of
observed=False is deprecated and will be changed to True in a future
version of pandas. Pass observed=False to retain current behavior or
observed=True to adopt the future default and silence this warning.
    survival_rate = df.groupby(['sex', 'class'])
['survived'].mean().reset_index();

```

1. Iris Flower Classification: Inbuilt Dataset : iris in Python

A.

- Display basic information and summary statistics of the dataset.
- Check for missing values in each column.

B. Create a scatter plot of sepal length vs. sepal width, colored by species.

```

df = sns.load_dataset('iris')

print(df.info())

print(df.describe())

sns.scatterplot(data=df, x='sepal_length', y='sepal_width',
hue='species')
plt.title('Sepal Length vs Sepal Width by Species')
plt.xlabel('Sepal Length')
plt.ylabel('Sepal Width')

```

```
plt.legend(title='Species')
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 150 entries, 0 to 149
```

```
Data columns (total 5 columns):
```


#	Column	Non-Null Count	Dtype
0	sepal_length	150 non-null	float64
1	sepal_width	150 non-null	float64
2	petal_length	150 non-null	float64
3	petal_width	150 non-null	float64
4	species	150 non-null	object

```
dtypes: float64(4), object(1)
```

```
memory usage: 6.0+ KB
```

```
None
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000



1. Distribution of Petal Length: Inbuilt dataset: iris in Python

Use histograms and density plots to visualize petal length distribution.


```
df = sns.load_dataset('iris')

# Histogram of petal length
sns.histplot(df['petal_length'], bins=15, color='skyblue')
plt.title('Histogram of Petal Length')
plt.xlabel('Petal Length')
plt.ylabel('Frequency')
plt.show()


# Density plot of petal length
sns.kdeplot(df['petal_length'], color='red')
plt.title('Density Plot of Petal Length')
plt.xlabel('Petal Length')
plt.ylabel('Density')
plt.show()

/Users/mac/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated
```

and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):



/Users/mac/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):



1. Ozone Levels Over Time: Inbuilt dataset: `seaborn.load_dataset('mpg')` in Python

A. find the number of unique car origins.

B. create a bar plot showing the average mpg for each origin.

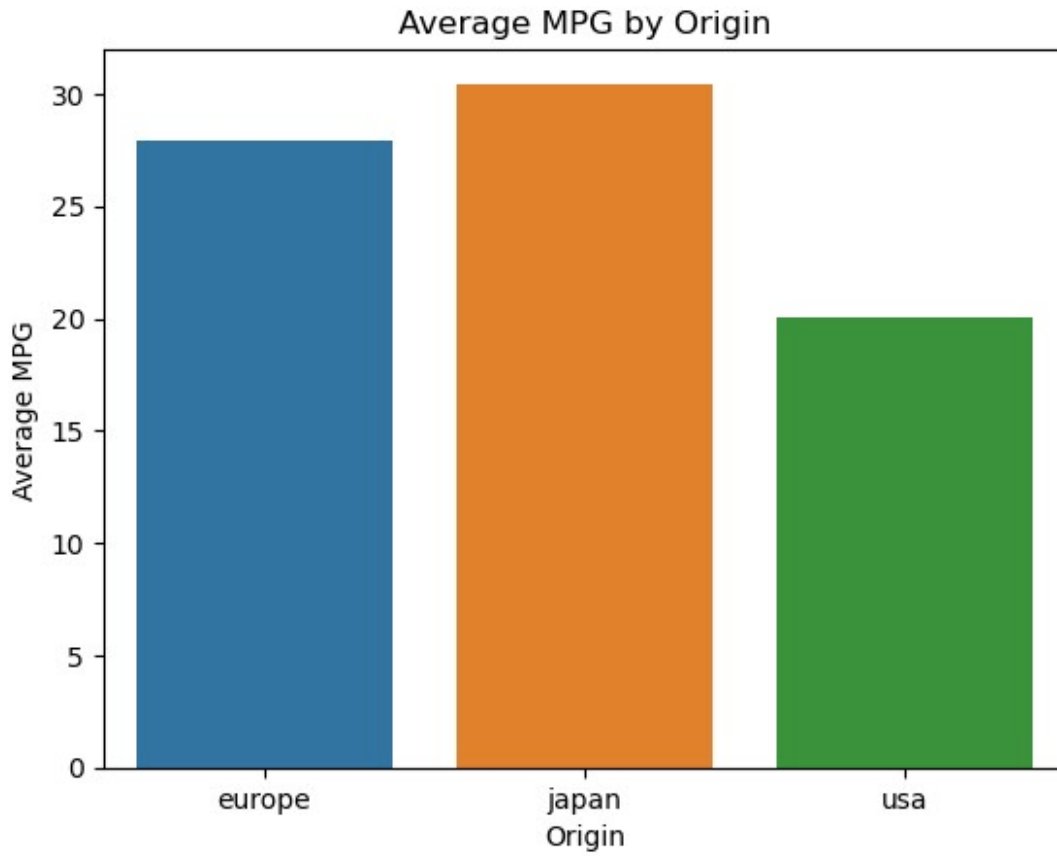
```
df = sns.load_dataset('mpg')

unique_origins = df['origin'].unique()
print("Unique origins:", unique_origins)

avg_mpg = df.groupby('origin')['mpg'].mean().reset_index()

sns.barplot(data=avg_mpg, x='origin', y='mpg')
plt.title('Average MPG by Origin')
plt.xlabel('Origin')
plt.ylabel('Average MPG')
plt.show()

Unique origins: ['usa' 'japan' 'europe']
```

1. Inbuilt dataset: `seaborn.load_dataset('diamonds')` in Python

A. Analyze how the average price of diamonds varies with the cut quality (e.g., Fair, Good, Ideal, etc.).

B. Create a box plot to visualize the distribution of diamond prices for each clarity level.

```
df = sns.load_dataset('diamonds')

# A. Average price by cut quality
avg_price_by_cut = df.groupby('cut')['price'].mean().reset_index()
print(avg_price_by_cut)

# B. Box plot: Price distribution by clarity
sns.boxplot(data=df, x='clarity', y='price', palette='coolwarm')
plt.title('Diamond Price Distribution by Clarity')
plt.xlabel('Clarity')
plt.ylabel('Price')
plt.show()
```

	cut	price
0	Ideal	3457.541970
1	Premium	4584.257704

```


2  Very Good  3981.759891
3      Good   3928.864452
4      Fair   4358.757764

```

```

/var/folders/kx/41v1tt6jlyx79h_8wk_hkl8w0000gn/T/
ipykernel_4467/1318471352.py:4: FutureWarning: The default of
observed=False is deprecated and will be changed to True in a future
version of pandas. Pass observed=False to retain current behavior or
observed=True to adopt the future default and silence this warning.
    avg_price_by_cut = df.groupby('cut')['price'].mean().reset_index()
/Users/mac/anaconda3/lib/python3.11/site-packages/seaborn/categorical.
py:641: FutureWarning: The default of observed=False is deprecated and
will be changed to True in a future version of pandas. Pass
observed=False to retain current behavior or observed=True to adopt
the future default and silence this warning.
    grouped_vals = vals.groupby(grouper)

```



1. A supermarket chain has collected sales data but has missing values and incorrect entries. The dataset is given below:

```
import pandas as pd
```

```
sales_data = pd.DataFrame({
```

```

"Transaction_ID": [101, 102, 103, 104],
>Date": pd.to_datetime(["2024-03-01", "2024-03-02", "2024-03-03",
"2024-03-04"]),
"Product": ["Apples", "Bread", "Milk", "Cheese"],
"Category": ["Fruits", "Bakery", "Dairy", "Dairy"],
"Quantity": [2, None, -1, 1],
"Price": [1.5, 2.0, 3.0, 5.0],
"Total_Sales": [3.0, None, -3.0, 5.0]

})

```

Write the code in Python for below problems

- Identify and handle missing values in Quantity and Total_Sales.
- Correct the incorrect Quantity values (negative values).
- Compute Total_Sales where missing.
- Summarize total sales per category.

```

sales_data = pd.DataFrame({
    "Transaction_ID": [101, 102, 103, 104],
    "Date": pd.to_datetime(["2024-03-01", "2024-03-02", "2024-03-03",
"2024-03-04"]),
    "Product": ["Apples", "Bread", "Milk", "Cheese"],
    "Category": ["Fruits", "Bakery", "Dairy", "Dairy"],
    "Quantity": [2, None, -1, 1],
    "Price": [1.5, 2.0, 3.0, 5.0],
    "Total_Sales": [3.0, None, -3.0, 5.0]
})

sales_data['Quantity'].fillna(0, inplace=True)
sales_data['Total_Sales'].fillna(0, inplace=True)

sales_data['Quantity'] = sales_data['Quantity'].apply(lambda x: abs(x)
if x < 0 else x)

sales_data['Total_Sales'] = sales_data['Quantity'] *
sales_data['Price']

category_sales = sales_data.groupby('Category')
['Total_Sales'].sum().reset_index()

print(sales_data)
print( category_sales)

```

	Transaction_ID	Date	Product	Category	Quantity	Price
Total_Sales						
0	101	2024-03-01	Apples	Fruits	2.0	1.5
3.0						
1	102	2024-03-02	Bread	Bakery	0.0	2.0
0.0						
2	103	2024-03-03	Milk	Dairy	1.0	3.0
3.0						
3	104	2024-03-04	Cheese	Dairy	1.0	5.0
5.0						
Category	Total_Sales					
0	Bakery	0.0				
1	Dairy	8.0				
2	Fruits	3.0				

1. Write the code in Python for below questions

```
import pandas as pd
```

```
df = pd.DataFrame({
```

```
    'Order_ID': [101, 102, 103, 103, 104, 105, 105],
    'Customer': ['Alice', 'Bob', None, None, 'Eve', 'Frank', 'Frank'],
    'Product': ['Laptop', 'Phone', 'Tablet', 'Tablet', 'Monitor', None,
                'Keyboard'],
    'Price': [1000, 500, 300, 300, 200, 150, 100],
    'Quantity': [2, None, 1, 1, 3, 2, 1]
```

```
})
```

** Identify and fill missing values:

- Fill missing Customer names with "Guest".
 - Fill missing Quantity values with the median quantity.
 - Fill missing Product values with "Unknown".
1. Remove duplicate Order_ID records, keeping the first occurrence
 2. Add a new column called "Total Amount" = Price * Quantity

```
df = pd.DataFrame({
    'Order_ID': [101, 102, 103, 103, 104, 105, 105],
    'Customer': ['Alice', 'Bob', None, None, 'Eve', 'Frank', 'Frank'],
    'Product': ['Laptop', 'Phone', 'Tablet', 'Tablet', 'Monitor',
                None, 'Keyboard'],
    'Price': [1000, 500, 300, 300, 200, 150, 100],
    'Quantity': [2, None, 1, 1, 3, 2, 1]
```

```

}))

# Fill missing values
df['Customer'].fillna('Guest', inplace=True)
df['Quantity'].fillna(df['Quantity'].median(), inplace=True)
df['Product'].fillna('Unknown', inplace=True)

# Remove duplicate Order_ID records, keeping the first
df_unique = df.drop_duplicates(subset='Order_ID', keep='first');

# Add "Total Amount" column
df_unique['Total Amount'] = df_unique['Price'] *
df_unique['Quantity'];
print(df_unique);

```

	Order_ID	Customer	Product	Price	Quantity	Total Amount
0	101	Alice	Laptop	1000	2.0	2000.0
1	102	Bob	Phone	500	1.5	750.0
2	103	Guest	Tablet	300	1.0	300.0
4	104	Eve	Monitor	200	3.0	600.0
5	105	Frank	Unknown	150	2.0	300.0

```

/var/folders/kx/41v1tt6jlyx79h_8wk_hkl8w0000gn/T/
ipykernel_4467/2724297360.py:18: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

df_unique['Total Amount'] = df_unique['Price'] *
df_unique['Quantity'];

```

1. Write the code in Python for below questions

```
df = pd.DataFrame({
```

```

'Transaction_ID': [1001, 1002, 1003, 1003, 1004, 1005],
'Customer': ['Alice', 'Bob', None, None, 'Eve', 'Frank'],
'Amount': [250, 400, None, 150, 700, 900],
'Discount': [10, 15, None, 5, None, 20]

```

```

})

```

1. Fill missing values:
 - Customer → "Guest"
 - Amount → mean of non-missing values

- Discount → replace None with 0
- Remove duplicate Transaction_IDs.
 - Add a new column "Final Amount", calculated as $\text{Amount} - (\text{Amount} * \text{Discount} / 100)$

```
df = pd.DataFrame({
    'Transaction_ID': [1001, 1002, 1003, 1003, 1004, 1005],
    'Customer': ['Alice', 'Bob', None, None, 'Eve', 'Frank'],
    'Amount': [250, 400, None, 150, 700, 900],
    'Discount': [10, 15, None, 5, None, 20]
})

# 1. Fill missing Customer with "Guest"
df['Customer'].fillna('Guest', inplace=True)

# 2. Fill missing Amount with the mean of non-missing values
mean_amount = df['Amount'].mean()
df['Amount'].fillna(mean_amount, inplace=True)

# 3. Replace missing Discount values with 0
df['Discount'].fillna(0, inplace=True)

# 4. Remove duplicate Transaction_IDs, keeping the first
df = df.drop_duplicates(subset='Transaction_ID', keep='first')

# 5. Add "Final Amount" = Amount - (Amount * Discount / 100)
df['Final Amount'] = df['Amount'] - (df['Amount'] * df['Discount'] / 100)

print(df)
```

	Transaction_ID	Customer	Amount	Discount	Final Amount
0	1001	Alice	250.0	10.0	225.0
1	1002	Bob	400.0	15.0	340.0
2	1003	Guest	480.0	0.0	480.0
4	1004	Eve	700.0	0.0	700.0
5	1005	Frank	900.0	20.0	720.0

- Write the code in Python for below questions

```
df = pd.DataFrame({
```

```
'Product_ID': [101, 102, 103, 103, 104, 105],
'Product_Name': ['Laptop', None, 'Tablet', 'Tablet', 'Monitor',
'Keyboard'],
'Stock': [50, None, 30, 30, 20, None],
'Price': [1000, 500, 300, 300, 200, 150]
```

```
})
```

1. Fill missing values:
 - Product_Name → "Unknown"
 - Stock → median of non-missing stock values
1. Remove duplicate Product_IDs.
2. Add a column "Stock Value", calculated as Stock * Price.

```
df = pd.DataFrame({
    'Product_ID': [101, 102, 103, 103, 104, 105],
    'Product_Name': ['Laptop', None, 'Tablet', 'Tablet', 'Monitor',
'Keyboard'],
    'Stock': [50, None, 30, 30, 20, None],
    'Price': [1000, 500, 300, 300, 200, 150]
})
```

```
# 1. Fill missing Product_Name with "Unknown"
```

```
df['Product_Name'].fillna('Unknown', inplace=True)
```

```
# 2. Fill missing Stock values with the median of non-missing stock
values
```

```
median_stock = df['Stock'].median()
```

```
df['Stock'].fillna(median_stock, inplace=True)
```

```
# 3. Remove duplicate Product_IDs, keeping the first
```

```
df = df.drop_duplicates(subset='Product_ID', keep='first')
```

```
# 4. Add "Stock Value" column = Stock * Price
```

```
df['Stock Value'] = df['Stock'] * df['Price']
```

```
print(df)
```

	Product_ID	Product_Name	Stock	Price	Stock Value
0	101	Laptop	50.0	1000	50000.0
1	102	Unknown	30.0	500	15000.0
2	103	Tablet	30.0	300	9000.0
4	104	Monitor	20.0	200	4000.0
5	105	Keyboard	30.0	150	4500.0

Golden question

1. Create a Python dataframe with at least 4 columns and 5 rows (you can generate a dataset of your choice). Perform the following tasks in Python :
 - Identify and handle missing values in the dataset.
 - Remove duplicate rows if any.
 - Add a new column based on existing data.
 - Generate at least two visualizations using Matplotlib or Seaborn to analyze trends or distributions in the dataset.

```

data = pd.DataFrame({
    'Customer': ['Alice', 'Bob', 'Charlie', 'Alice', np.nan],
    'Product': ['Laptop', 'Phone', 'Tablet', 'Laptop', 'Phone'],
    'Quantity': [1, 2, np.nan, 1, 2],
    'Price': [1000, 500, 300, 1000, 500]
})

print(data)

data['Customer'].fillna('Guest', inplace=True)
data['Quantity'].fillna(data['Quantity'].median(), inplace=True)

data.drop_duplicates(inplace=True)

data['Total'] = data['Quantity'] * data['Price']
print(data)

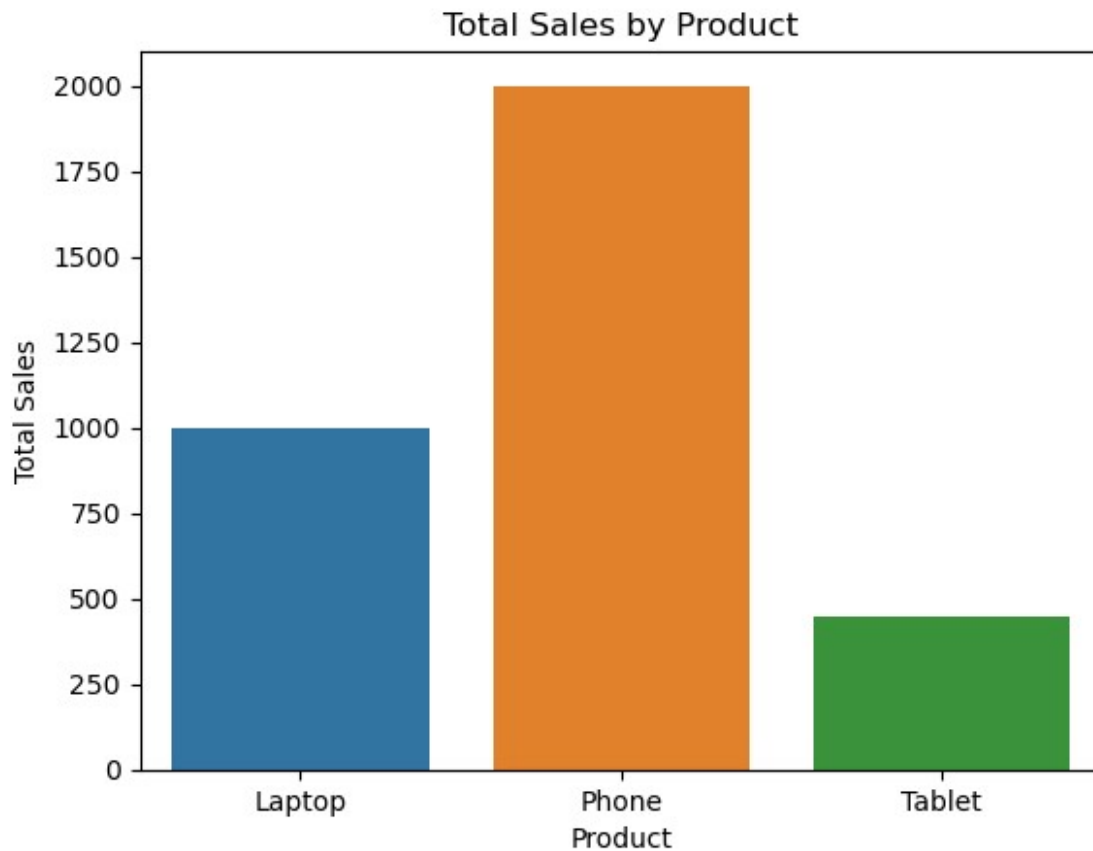
sns.barplot(data=data, x='Product', y='Total', estimator=sum)
plt.title("Total Sales by Product")
plt.ylabel("Total Sales")
plt.xlabel("Product")
plt.show()

sns.histplot(data['Quantity'], bins=5)
plt.title("Distribution of Quantity Purchased")
plt.xlabel("Quantity")
plt.ylabel("Frequency")
plt.show()


```

	Customer	Product	Quantity	Price
0	Alice	Laptop	1.0	1000
1	Bob	Phone	2.0	500
2	Charlie	Tablet	NaN	300
3	Alice	Laptop	1.0	1000
4	NaN	Phone	2.0	500

	Customer	Product	Quantity	Price	Total
0	Alice	Laptop	1.0	1000	1000.0
1	Bob	Phone	2.0	500	1000.0
2	Charlie	Tablet	1.5	300	450.0
4	Guest	Phone	2.0	500	1000.0



```
/Users/mac/anaconda3/lib/python3.11/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
  with pd.option_context('mode.use_inf_as_na', True):
```



```
# 1. Air Quality Analysis: Inbuilt dataset: airquality in R
# A. Filter the records for the month of July.
# B. Group the data by Month and calculate the average Ozone.
# C. Use a pipe operator to fetch records where Ozone > 50.
```

```
library(dplyr)
```

```
data("airquality")
```

```
# A. Filter records for July (Month = 7)
```

```
july_data <- airquality %>%
  filter(Month == 7)
```

```
print(july_data)
```

```
# B. Group by Month and calculate average Ozone
```

```
ozone_avg <- airquality %>%
  group_by(Month) %>%
  summarise(Avg_Ozone = mean(Ozone), na.rm = TRUE)
```

```
print(ozone_avg)
```

```
# C. Use pipe to fetch records with Ozone > 50
```

```
high_ozone <- airquality %>%
  filter(Ozone > 50)
```

```
print(high_ozone)
```

```
# 3. Car Performance Analysis: Inbuilt dataset: mtcars in R
```

```
# A. Compare the fuel efficiency (mpg) of automatic vs. manual transmission cars.
```

```
# B. Identify the relationship between horsepower (hp) and fuel consumption.
```

```
library(dplyr)
```

```
library(ggplot2)
```

```
# Add a readable label for transmission
```

```
mtcars$Transmission <- ifelse(mtcars$am == 0, "Automatic", "Manual")
```

```
# Calculate average mpg by transmission
```

```
avg_mpg <- mtcars %>%
  group_by(Transmission) %>%
  summarise(Average_MPG = mean(mpg))
```

```
print(avg_mpg)
```

```
# Bar plot for comparison
```

```
ggplot(avg_mpg, aes(x = Transmission, y = Average_MPG, fill = Transmission)) +
  geom_bar(stat = "identity") +
  labs(title = "Fuel Efficiency by Transmission Type",
       x = "Transmission Type",
```

```

    y = "Average MPG") +
theme_minimal()

# 5. Titanic Survival Analysis: Inbuilt Dataset: Titanic in R

# A. Compute the total number of passengers by gender and class.

# B. Calculate the percentage of passengers who survived, grouped by class.

library(titanic)
library(dplyr)

data <- titanic_train

# A. Total number of passengers by gender and class
passenger_counts <- data %>%
  group_by(Sex, Pclass) %>%
  summarise(Total_Passengers = n())

print(passenger_counts)

# B. Percentage of passengers who survived, grouped by class
survival_by_class <- data %>%
  group_by(Pclass) %>%
  summarise(Survival_Rate = mean(Survived) * 100)

print(survival_by_class)

# 5. Dataset: PlantGrowth (inbuilt in R)

# A. Compute the average weight of plants in each treatment group.
# B. Create a bar chart to visualize the average plant weights per group.

library(dplyr)
library(ggplot2)

data("PlantGrowth")

# A. Compute average weight by group
avg_weight <- PlantGrowth %>%
  group_by(group) %>%
  summarise(Avg_Weight = mean(weight))

print(avg_weight)

# B. Bar chart of average weight per group
ggplot(avg_weight, aes(x = group, y = Avg_Weight, fill = group)) +
  geom_bar(stat = "identity") +
  labs(title = "Average Plant Weight by Group",
       x = "Treatment Group",

```

```

    y = "Average Weight") +
  theme_minimal()

# 7. Iris Flower Classification: Inbuilt Dataset : iris in R

# A. Calculate the average petal length and petal width for each species.
# B. Create a scatter plot of Sepal.Length vs Sepal.Width colored by species

library(dplyr)
library(ggplot2)

data("iris")

# A. Average Petal.Length and Petal.Width by Species
avg_petal <- iris %>%
  group_by(Species) %>%
  summarise(
    Avg_Petal_Length = mean(Petal.Length),
    Avg_Petal_Width = mean(Petal.Width)
  )

print(avg_petal)

# B. Scatter plot of Sepal.Length vs Sepal.Width by Species
ggplot(iris, aes(x = Sepal.Length, y = Sepal.Width, color = Species)) +
  geom_point(size = 3) +
  labs(title = "Sepal Dimensions by Species",
       x = "Sepal Length",
       y = "Sepal Width") +
  theme_minimal()

# 9. Distribution of Petal Length: Inbuilt dataset: iris in R

# Use histograms and density plots to visualize petal length distribution.

library(ggplot2)

data("iris")

# Histogram of Petal Length
ggplot(iris, aes(x = Petal.Length)) +
  geom_histogram(binwidth = 0.5, fill = "skyblue", color = "black") +
  labs(title = "Histogram of Petal Length",
       x = "Petal Length",
       y = "Frequency") +
  theme_minimal()

# Density plot of Petal Length
ggplot(iris, aes(x = Petal.Length)) +
  geom_density(fill = "lightgreen", alpha = 0.6) +
  labs(title = "Density Plot of Petal Length",
       x = "Petal Length",

```

```

    y = "Density") +
theme_minimal()

# 11. Dataset: mtcars (inbuilt in R)
# A. Filter and show details of cars with horsepower (hp) greater than 150.
# B. Create a scatter plot showing the relationship between horsepower (hp) and fuel efficiency

library(ggplot2)
library(dplyr)

# Load dataset
data("mtcars")

high_hp_cars <- mtcars %>% filter(hp > 150)

print(high_hp_cars)

# Scatter plot
ggplot(mtcars, aes(x = hp, y = mpg)) +
  geom_point(color = "steelblue", size = 3) +
  labs(title = "Horsepower vs. Fuel Efficiency",
       x = "Horsepower (hp)",
       y = "Miles per Gallon (mpg)") +
  theme_minimal()

# 13. CO2 Emissions : Inbuilt dataset: CO2 in R

# A. Compare CO2 uptake between different treatment groups.
# B. Analyze which factors significantly affect CO2 levels.

library(dplyr)
library(ggplot2)

data("CO2")

# A. Average CO2 uptake by Treatment group
avg_uptake <- CO2 %>%
  group_by(Treatment) %>%
  summarise(Avg_Uptake = mean(uptake))

print(avg_uptake)

# B. Scatter plot: CO2 uptake vs. concentration, colored by Plant Type
ggplot(CO2, aes(x = conc, y = uptake, color = Type)) +
  geom_point(size = 3) +
  labs(title = "CO2 Uptake by Concentration and Plant Type",
       x = "CO2 Concentration (ppm)",
       y = "CO2 Uptake",
       color = "Plant Type") +
  theme_minimal()

```

```

# 15. A supermarket chain has collected sales data but has missing values and incorrect entries

# sales_data <- data.frame(

#   Transaction_ID = c(101, 102, 103, 104),

#   Date = as.Date(c("2024-03-01", "2024-03-02", "2024-03-03", "2024-03-04")),

#   Product = c("Apples", "Bread", "Milk", "Cheese"),

#   Category = c("Fruits", "Bakery", "Dairy", "Dairy"),

#   Quantity = c(2, NA, -1, 1),

#   Price = c(1.5, 2.0, 3.0, 5.0),

#   Total_Sales = c(3.0, NA, -3.0, 5.0)

# )

# Write the code in R for below problems:

# Identify and handle missing values in Quantity and Total_Sales.
# Correct the incorrect Quantity values (negative values).
# Compute Total_Sales where missing.
# Summarize total sales per category.

sales_data <- data.frame(
  Transaction_ID = c(101, 102, 103, 104),
  Date = as.Date(c("2024-03-01", "2024-03-02", "2024-03-03", "2024-03-04")),
  Product = c("Apples", "Bread", "Milk", "Cheese"),
  Category = c("Fruits", "Bakery", "Dairy", "Dairy"),
  Quantity = c(2, NA, -1, 1),
  Price = c(1.5, 2.0, 3.0, 5.0),
  Total_Sales = c(3.0, NA, -3.0, 5.0)
)

# 1. Handle missing values in Quantity and Total_Sales
# Replace missing Quantity with the median
sales_data$Quantity[is.na(sales_data$Quantity)] <- median(sales_data$Quantity, na.rm = TRUE)

# Replace missing Total_Sales with 0
sales_data$Total_Sales[is.na(sales_data$Total_Sales)] <- 0

# 2. Correct negative Quantity values
sales_data$Quantity[sales_data$Quantity < 0] <- abs(sales_data$Quantity[sales_data$Quantity < 0])

# 3. Recompute Total_Sales where it's 0 or wrong
sales_data$Total_Sales <- sales_data$Quantity * sales_data$Price

# 4. Summarize total sales per category
library(dplyr)

```

```

category_summary <- sales_data %>%
  group_by(Category) %>%
  summarise(Total_Sales_Sum = sum(Total_Sales))

print(category_summary)

# Golden Question

# 2. Using any built-in dataset in R, perform the following tasks:

# Data Manipulation using dplyr:

# Select relevant columns for analysis.
# Filter the dataset based on a meaningful condition.
# Create a new derived column using existing data.
# Group the data and compute summary statistics.
# Arrange the dataset meaningfully (e.g., in ascending or descending order).
# Data Visualization using ggplot2:

# Create at least two visualizations to explore trends or distributions in the dataset
# Use appropriate aesthetics such as color, size, and facets.
# Add clear axis labels, a title, and a legend where necessary.

library(dplyr)
library(ggplot2)

head(mtcars)

# Data Manipulation
manipulated_data <- mtcars %>%
  select(mpg, cyl, hp, gear) %>%
  filter(hp > 100) %>%
  mutate(Efficiency = mpg / cyl) %>%
  group_by(gear) %>%
  summarise(
    Avg_MPG = mean(mpg),
    Avg_HP = mean(hp),
    Count = n()
  ) %>%
  arrange(desc(Avg_MPG))

print(manipulated_data)

# Scatter Plot - HP vs MPG
ggplot(mtcars, aes(x = hp, y = mpg)) +
  geom_point(size = 3) +
  labs(
    title = "Horsepower vs MPG",
    x = "Horsepower (hp)",

```



```
  y = "Miles Per Gallon (mpg)",
  color = "Cylinders"
) +
theme_minimal()

# Boxplot - MPG by Gear
ggplot(mtcars, aes(x = factor(gear), y = mpg)) +
  geom_boxplot() +
  labs(
    title = "Distribution of MPG by Number of Gears",
    x = "Number of Gears",
    y = "Miles Per Gallon (mpg)"
  ) +
  theme_minimal()
```